

VibeHaus - AI Responsibility Statement

Author: Sakshi · Capstone ·

VibeHaus uses AI in three places: **Vibe Match** (recommending compatible people), **Icebreakers** (suggested conversation starters), and **Mr. Vibe** (an activity-based search assistant, Stretch). This statement covers data, human review, limitations, and escalation.

What data goes into AI (and what doesn't)

- **In:** user-provided, non-sensitive signals — interests, social energy, intent, city, age band, and in-session like/save feedback.
- **Not in / handled carefully:** optional profile fields (languages, hometown/country, new-in-town) are self-declared, optional, and never required; they are not used as hard filters and not used to match on protected attributes. **No sensitive/protected attributes (e.g. ethnicity) are collected.** No payment or ID data is processed in the MVP.

How matching works (transparency over black box)

Compatibility is **rule/attribute-based** — the app weights shared interests, social energy, intent, availability, and connection preference. The **LLM's job is to generate the plain-language "why you might vibe" explanation, Icebreakers, and Mr. Vibe's conversational replies**, not to compute the match or search result itself. There is no opaque personality profiling and no model training on user data in the MVP: like/save may influence which profiles show next within the product, but user interactions are not used to train external AI models. Every recommendation is shown with a **banded score + reason listing every genuinely shared interest** ("Strong Vibe — both new to Berlin, into photography, prefer smaller groups"), never an unexplained precise percentage.

Mr. Vibe: a design decision made honest by testing

Mr. Vibe was originally designed to accept free-text descriptions of what a user wanted to do, interpreted by a constrained LLM. During testing, this free-text path proved unreliable — it sometimes searched for or displayed a different activity than the one typed. Rather than ship an unpredictable live-parsing feature, **free-text input was removed and the chip-based options expanded to cover the app's complete interest list**, so every request is unambiguous and every result is genuinely searched against what was actually selected. Mr. Vibe's conversation flow (greetings, follow-up questions, fallback messages) is scripted for reliability and cost; only the underlying compatibility search reuses the app's existing rule-based engine. Mr. Vibe **never performs real account actions** (deleting an account, blocking, changing settings) regardless of

how it is asked, and if a user's message signals emotional distress (e.g. loneliness), it responds with a single brief, warm, scripted acknowledgment before redirecting to its actual job — it does not attempt counseling or open-ended emotional support.

Human review & control

The user is always in the loop: they can ignore recommendations, browse freely, edit preferences anytime, and react with like/save/scroll-past — all private, with no notification to the other person. Liking someone only opens a further choice (**Send Vibemate Request** or **Maybe Later**, which saves to Saved Vibes instead); nothing is sent automatically. A **Vibemate** relationship — and any contact at all — requires **mutual acceptance**, auto-adding both people to **My Haus** and unlocking Chat. Coffee Vibe and Activity Vibe are separate, further-consent suggestions within that relationship, each needing its own accept, and either party can cancel a confirmed one afterward. Icebreakers are optional and editable — suggestions, never sent automatically. This graduated model (private signal → mutual request → relationship → engagement mode) reflects one principle: **increasing intimacy requires increasing consent**.

Known limitations & user-facing failure modes

- The AI can suggest poor matches, especially with sparse profiles or a thin seeded pool → framed clearly as *suggestions*, with a filter/chip fallback and best-available handling, never an empty or misleadingly confident screen.
- Explained-match reasons are generated from profile overlap and could occasionally read generically → kept short and grounded in real shared attributes.
- Small, expat-skewed discovery data means matching assumptions are unvalidated at scale → stated openly; broader testing is the next step.
- Mr. Vibe's chip-based design trades some input flexibility for reliability — a deliberate choice after testing surfaced the free-text problem above, not an oversight.

Escalation (mistakes, abuse, edge cases)

Uncertain or failed AI output falls back to plain, filter/chip-based discovery rather than guessing. Because sensitive fields are optional and never used as filters, the app avoids enabling protected-attribute screening of any kind. **Block, Report, and Remove Vibemate are built into the MVP**, not deferred — removing a Vibemate immediately and consistently revokes their access to Chat, Coffee Vibe, and Activity Vibe everywhere their profile can be viewed, checked live rather than cached. Identity verification and automated moderation remain roadmap items (see backlog, Next/Later tiers), and a fuller safety review is required before any real-world launch.

One-line summary

VibeHaus's AI is a *transparent, optional assistant* that explains itself and keeps the human in control — chosen deliberately over an opaque compatibility score or an unpredictable free-text assistant, because trust matters more than magic for a friendship product.